

# etl

September 28, 2021

## 1 ETL Processes

Use this notebook to develop the ETL process for each of your tables before completing the `etl.py` file to load the whole datasets.

```
In [1]: import os
import glob
import psycopg2
import pandas as pd
from sql_queries import import *
import json
```

```
In [2]: conn = psycopg2.connect("host=127.0.0.1 dbname=sparkifydb user=student password=student")
cur = conn.cursor()
```

```
In [3]: def get_files(filepath):
    all_files = []
    for root, dirs, files in os.walk(filepath):
        files = glob.glob(os.path.join(root, '*.json'))
        for f in files :
            all_files.append(os.path.abspath(f))

    return all_files
```

## 2 Process song\_data

In this first part, you'll perform ETL on the first dataset, `song_data`, to create the songs and artists dimensional tables.

Let's perform ETL on a single song file and load a single record into each table to start. - Use the `get_files` function provided above to get a list of all song JSON files in `data/song_data` - Select the first song in this list - Read the song file and view the data

```
In [4]: song_files = get_files('data/song_data')
```

```
In [5]: filepath = song_files[0]
print(filepath)
```

```
/home/workspace/data/song_data/A/A/A/TRAAAAW128F429D538.json
```

```
In [6]: with open(filepath, 'r') as f: # getting data as dict (load_json not working due to scalar)
        data = json.load(f)

        df = pd.DataFrame({0: data}) # manually indexing since data is scalar
        df = df.transpose() # transposing for column view
        df.head()
```

```
Out[6]:
```

	artist_id	artist_latitude	artist_location	artist_longitude	
0	ARD7TVE1187B99BFB1	None	California - LA	None	

  

	artist_name	duration	num_songs	song_id	title	year
0	Casual	218.932	1	SOMZWCG12A8C13C480	I Didn't Mean To	0

## 2.1 #1: songs Table

### Extract Data for Songs Table

- Select columns for song ID, title, artist ID, year, and duration
- Use `df.values` to select just the values from the dataframe
- Index to select the first (only) record in the dataframe
- Convert the array to a list and set it to `song_data`

```
In [7]: song_list = [] # list of dictionaries to host song data to be converted to pandas df

        for i in range(len(song_files)): # iterating through list of song files
            path = song_files[i]
            with open(path, 'r') as f: # appending each to the list created above
                song_list.append(json.load(f))

        print(len(song_list))
        song_df = pd.DataFrame(song_list) # converting list of dictionaries to pandas df
        print(song_df.dtypes)
        print(song_df.shape)
        song_df.head()
```

```
81
artist_id      object
artist_latitude float64
artist_location object
artist_longitude float64
artist_name     object
duration        float64
num_songs       int64
song_id         object
title           object
year            int64
```

```
dtype: object
(81, 10)
```

```
Out[7]:
```

	artist_id	artist_latitude	artist_location	\
0	ARD7TVE1187B99BFB1	NaN	California - LA	
1	ARNTLGG11E2835DDB9	NaN		
2	AR8ZCNI1187B9A069B	NaN		
3	AR10USD1187B99F3F1	NaN	Burlington, Ontario, Canada	
4	ARMJAGH1187FB546F3	35.14968	Memphis, TN	

	artist_longitude	artist_name	duration	num_songs	\
0	NaN	Casual	218.93179	1	
1	NaN	Clp	266.39628	1	
2	NaN	Planet P Project	269.81832	1	
3	NaN	Tweeterfriendly Music	189.57016	1	
4	-90.04892	The Box Tops	148.03546	1	

	song_id	title	year
0	SOMZWCG12A8C13C480	I Didn't Mean To	0
1	SOUDSGM12AC9618304	Insatiable (Instrumental Version)	0
2	SOIAZJW12AB01853F1	Pink World	1984
3	SOHKNRJ12A6701D1F8	Drop of Rain	0
4	SOCIWDW12A8C13D406	Soul Deep	1969

```
In [8]: song_data_subset = song_df[['song_id', 'title', 'artist_id', 'year', 'duration']]
song_data_array = song_data_subset.values
song_data_list = song_data_array.tolist()
first_song_array = song_data_array[0]
song_data = first_song_array.tolist()
```

```
In [9]: #for i in range(len(song_data_list)):
#       print(song_data_list[i])
# for test purposes only
```

**Insert Record into Song Table** Implement the `song_table_insert` query in `sql_queries.py` and run the cell below to insert a record for this song into the songs table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the songs table in the sparkify database.

```
In [10]: #cur.execute(song_table_insert, first_song_tuple) for SINGLE row as requested above
#cur.execute(song_table_insert, song_data)

#conn.commit()

# insert all data (which is needed for the final implementation)

for i in range(len(song_data_list)):
```

```

current_song = song_data_list[i]
cur.execute(song_table_insert, current_song)
conn.commit()

```

Run `test.ipynb` to see if you've successfully added a record to this table.

## 2.2 #2: artists Table

### Extract Data for Artists Table

- Select columns for artist ID, name, location, latitude, and longitude
- Use `df.values` to select just the values from the dataframe
- Index to select the first (only) record in the dataframe
- Convert the array to a list and set it to `artist_data`

```

In [11]: artist_data = song_df[['artist_id', 'artist_name', 'artist_location', 'artist_latitude',
artist_array = artist_data.values
artist_list = artist_array.tolist()
first_artist_array = artist_array[0]
artist_data = first_artist_array.tolist()
artist_data

```

```

Out[11]: ['ARD7TVE1187B99BFB1', 'Casual', 'California - LA', nan, nan]

```

**Insert Record into Artist Table** Implement the `artist_table_insert` query in `sql_queries.py` and run the cell below to insert a record for this song's artist into the `artists` table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the `artists` table in the `sparkify` database.

```

In [12]: # insert single (as requested in the question...)
#cur.execute(artist_table_insert, artist_data)
#conn.commit()

# insert all data (which is needed for the final implementation)

for i in range(len(artist_list)):
    current_artist = artist_list[i]
    cur.execute(artist_table_insert, current_artist)
    conn.commit()

```

Run `test.ipynb` to see if you've successfully added a record to this table.

## 3 Process log\_data

In this part, you'll perform ETL on the second dataset, `log_data`, to create the time and users dimensional tables, as well as the `songplays` fact table.

Let's perform ETL on a single log file and load a single record into each table. - Use the `get_files` function provided above to get a list of all log JSON files in `data/log_data` - Select the first log file in this list - Read the log file and view the data

```
In [13]: log_files = get_files('data/log_data')
```

```
In [14]: filepath = log_files[0]
         print(filepath)
```

```
/home/workspace/data/log_data/2018/11/2018-11-30-events.json
```

```
In [15]: df = pd.read_json (filepath, lines = True)
         print(df.shape)
         df.head()
```

```
(388, 18)
```

```
Out[15]:
```

	artist	auth	firstName	gender	itemInSession	lastName	\
0	Stephen Lynch	Logged In	Jayden	M	0	Bell	
1	Manowar	Logged In	Jacob	M	0	Klein	
2	Morcheeba	Logged In	Jacob	M	1	Klein	
3	Maroon 5	Logged In	Jacob	M	2	Klein	
4	Train	Logged In	Jacob	M	3	Klein	

  

	length	level	location	method	page	\
0	182.85669	free	Dallas-Fort Worth-Arlington, TX	PUT	NextSong	
1	247.56200	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
2	257.41016	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
3	231.23546	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
4	216.76363	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	

  

	registration	sessionId	song	status	\
0	1.540992e+12	829	Jim Henson's Dead	200	
1	1.540558e+12	1049	Shell Shock	200	
2	1.540558e+12	1049	Women Lose Weight (Feat: Slick Rick)	200	
3	1.540558e+12	1049	Won't Go Home Without You	200	
4	1.540558e+12	1049	Hey_ Soul Sister	200	

  

	ts	userAgent	userId
0	1543537327796	Mozilla/5.0 (compatible; MSIE 10.0; Windows NT...	91
1	1543540121796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
2	1543540368796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
3	1543540625796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
4	1543540856796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73

### 3.1 #3: time Table

#### Extract Data for Time Table

- Filter records by NextSong action
- Convert the ts timestamp column to datetime

- Hint: the current timestamp is in milliseconds
- Extract the timestamp, hour, day, week of year, month, year, and weekday from the `ts` column and set `time_data` to a list containing these values in order
- Hint: use pandas' `dt attribute` to access easily datetimelike properties.
- Specify labels for these columns and set to `column_labels`
- Create a dataframe, `time_df`, containing the time data for this file by combining `column_labels` and `time_data` into a dictionary and converting this into a dataframe

```
In [16]: df_ns = df[df['page'] == 'NextSong']
         print(df_ns.shape)
         df_ns.head()
```

(330, 18)

```
Out[16]:
```

	artist	auth	firstName	gender	itemInSession	lastName	\
0	Stephen Lynch	Logged In	Jayden	M	0	Bell	
1	Manowar	Logged In	Jacob	M	0	Klein	
2	Morcheeba	Logged In	Jacob	M	1	Klein	
3	Maroon 5	Logged In	Jacob	M	2	Klein	
4	Train	Logged In	Jacob	M	3	Klein	

  

	length	level	location	method	page	\
0	182.85669	free	Dallas-Fort Worth-Arlington, TX	PUT	NextSong	
1	247.56200	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
2	257.41016	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
3	231.23546	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
4	216.76363	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	

  

	registration	sessionId	song	status	\
0	1.540992e+12	829	Jim Henson's Dead	200	
1	1.540558e+12	1049	Shell Shock	200	
2	1.540558e+12	1049	Women Lose Weight (Feat: Slick Rick)	200	
3	1.540558e+12	1049	Won't Go Home Without You	200	
4	1.540558e+12	1049	Hey_ Soul Sister	200	

  

	ts	userAgent	userId
0	1543537327796	Mozilla/5.0 (compatible; MSIE 10.0; Windows NT...	91
1	1543540121796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
2	1543540368796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
3	1543540625796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73
4	1543540856796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	73

```
In [17]: t = df_ns.copy()
         t['ts'] = pd.to_datetime(df_ns['ts'], unit='ms')
         t.head()
```

```
Out[17]:
```

	artist	auth	firstName	gender	itemInSession	lastName	\
0	Stephen Lynch	Logged In	Jayden	M	0	Bell	

1	Manowar	Logged In	Jacob	M	0	Klein
2	Morcheeba	Logged In	Jacob	M	1	Klein
3	Maroon 5	Logged In	Jacob	M	2	Klein
4	Train	Logged In	Jacob	M	3	Klein

	length	level	location	method	page	\
0	182.85669	free	Dallas-Fort Worth-Arlington, TX	PUT	NextSong	
1	247.56200	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
2	257.41016	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
3	231.23546	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	
4	216.76363	paid	Tampa-St. Petersburg-Clearwater, FL	PUT	NextSong	

	registration	sessionId	song	status	\
0	1.540992e+12	829	Jim Henson's Dead	200	
1	1.540558e+12	1049	Shell Shock	200	
2	1.540558e+12	1049	Women Lose Weight (Feat: Slick Rick)	200	
3	1.540558e+12	1049	Won't Go Home Without You	200	
4	1.540558e+12	1049	Hey_ Soul Sister	200	

	ts	userAgent	\
0	2018-11-30 00:22:07.796	Mozilla/5.0 (compatible; MSIE 10.0; Windows NT...	
1	2018-11-30 01:08:41.796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	
2	2018-11-30 01:12:48.796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	
3	2018-11-30 01:17:05.796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	
4	2018-11-30 01:20:56.796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	

	userId
0	91
1	73
2	73
3	73
4	73

```
In [18]: t['year'] = pd.DatetimeIndex(t['ts']).year
t['month'] = pd.DatetimeIndex(t['ts']).month
t['week'] = pd.DatetimeIndex(t['ts']).week
t['day'] = pd.DatetimeIndex(t['ts']).day
t['weekday'] = pd.DatetimeIndex(t['ts']).weekday
t['hour'] = pd.DatetimeIndex(t['ts']).hour
```

```
# converting df to dict
df_dict = t.T.to_dict()
```

```
list_of_time_data = []
```

```
# only want dates
for key, value in df_dict.items():
```

```

        entry = value
        dates_dict = {k: v for k, v in entry.items() if k in ['year', 'month', 'week', 'day']}
        list_of_time_data.append(dates_dict)

    for i in range(len(list_of_time_data)):
        print(list_of_time_data[i])
        if i > 5: break # keeping the view clean

{'ts': Timestamp('2018-11-30 00:22:07.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:08:41.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:12:48.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:17:05.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:20:56.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:24:32.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}
{'ts': Timestamp('2018-11-30 01:28:19.796000'), 'year': 2018, 'month': 11, 'week': 48, 'day': 30}

```

```

In [19]: time_df = pd.DataFrame(list_of_time_data)

# want order consistent with SQL INSERT query
time_df = time_df[['ts', 'hour', 'day', 'week', 'month', 'year', 'weekday']]

time_df.head()

```

```

Out[19]:
   ts      hour  day  week  month  year  weekday
0 2018-11-30 00:22:07.796    0   30   48    11  2018         4
1 2018-11-30 01:08:41.796    1   30   48    11  2018         4
2 2018-11-30 01:12:48.796    1   30   48    11  2018         4
3 2018-11-30 01:17:05.796    1   30   48    11  2018         4
4 2018-11-30 01:20:56.796    1   30   48    11  2018         4

```

**Insert Records into Time Table** Implement the `time_table_insert` query in `sql_queries.py` and run the cell below to insert records for the timestamps in this log file into the time table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the time table in the sparkify database.

```

In [20]: for i, row in time_df.iterrows():
        cur.execute(time_table_insert, list(row))
        conn.commit()

```

Run `test.ipynb` to see if you've successfully added records to this table.

## 3.2 #4: users Table

### Extract Data for Users Table

- Select columns for user ID, first name, last name, gender and level and set to `user_df`

```

In [21]: user_df = df[['userId', 'firstName', 'lastName', 'gender', 'level']]
        user_df.head()

```



```
Out[21]:  userId  firstName  lastName  gender  level
         0      91    Jayden    Bell      M    free
         1      73     Jacob    Klein     M    paid
         2      73     Jacob    Klein     M    paid
         3      73     Jacob    Klein     M    paid
         4      73     Jacob    Klein     M    paid
```

**Insert Records into Users Table** Implement the `user_table_insert` query in `sql_queries.py` and run the cell below to insert records for the users in this log file into the `users` table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the `users` table in the sparkify database.

```
In [22]: for i, row in user_df.iterrows():
         cur.execute(user_table_insert, row)
         conn.commit()
```

Run `test.ipynb` to see if you've successfully added records to this table.

### 3.3 #5: songplays Table

**Extract Data and Songplays Table** This one is a little more complicated since information from the `songs` table, `artists` table, and original log file are all needed for the `songplays` table. Since the log file does not specify an ID for either the song or the artist, you'll need to get the song ID and artist ID by querying the `songs` and `artists` tables to find matches based on song title, artist name, and song duration time. - Implement the `song_select` query in `sql_queries.py` to find the song ID and artist ID based on the title, artist name, and duration of a song. - Select the timestamp, user ID, level, song ID, artist ID, session ID, location, and user agent and set to `songplay_data`

#### Insert Records into Songplays Table

- Implement the `songplay_table_insert` query and run the cell below to insert records for the songplay actions in this log file into the `songplays` table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the `songplays` table in the sparkify database.

```
In [23]: print(df.dtypes)
         print(df.shape)
         df.head()
```

```
artist          object
auth            object
firstName       object
gender          object
itemInSession   int64
lastName        object
length          float64
level           object
location        object
```

```

method      object
page        object
registration float64
sessionId   int64
song        object
status      int64
ts          int64
userAgent   object
userId      object
dtype: object
(388, 18)

```

```

Out[23]:
      artist  auth firstName gender itemInSession lastName \
0  Stephen Lynch Logged In   Jayden      M           0      Bell
1    Manowar Logged In    Jacob      M           0     Klein
2  Morcheeba Logged In    Jacob      M           1     Klein
3  Maroon 5 Logged In    Jacob      M           2     Klein
4    Train Logged In    Jacob      M           3     Klein

      length level                location method  page \
0  182.85669  free  Dallas-Fort Worth-Arlington, TX  PUT  NextSong
1  247.56200  paid  Tampa-St. Petersburg-Clearwater, FL  PUT  NextSong
2  257.41016  paid  Tampa-St. Petersburg-Clearwater, FL  PUT  NextSong
3  231.23546  paid  Tampa-St. Petersburg-Clearwater, FL  PUT  NextSong
4  216.76363  paid  Tampa-St. Petersburg-Clearwater, FL  PUT  NextSong

      registration  sessionId                song  status \
0  1.540992e+12      829  Jim Henson's Dead      200
1  1.540558e+12     1049  Shell Shock      200
2  1.540558e+12     1049  Women Lose Weight (Feat: Slick Rick) 200
3  1.540558e+12     1049  Won't Go Home Without You      200
4  1.540558e+12     1049  Hey_ Soul Sister      200

      ts                userAgent  userId
0  1543537327796  Mozilla/5.0 (compatible; MSIE 10.0; Windows NT...    91
1  1543540121796  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    73
2  1543540368796  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    73
3  1543540625796  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    73
4  1543540856796  "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...    73

```

```

In [24]: for index, row in df.iterrows():
          cur.execute(song_select, (row.song, row.artist, row.length)) # running the songselect query
          results = cur.fetchone()
          if results:
              print(results)
              songid, artistid = results

```

```
else:
    songid, artistid = None, None
```

Run `test.ipynb` to see if you've successfully added records to this table.

## 4 Close Connection to Sparkify Database

```
In [25]: conn.close()
```

## 5 Implement `etl.py`

Use what you've completed in this notebook to implement `etl.py`.

```
In [26]: %run etl.py
```

```
81 files found in data/song_data
1/81 files processed.
2/81 files processed.
3/81 files processed.
4/81 files processed.
5/81 files processed.
6/81 files processed.
7/81 files processed.
8/81 files processed.
9/81 files processed.
10/81 files processed.
11/81 files processed.
12/81 files processed.
13/81 files processed.
14/81 files processed.
15/81 files processed.
16/81 files processed.
17/81 files processed.
18/81 files processed.
19/81 files processed.
20/81 files processed.
21/81 files processed.
22/81 files processed.
23/81 files processed.
24/81 files processed.
25/81 files processed.
26/81 files processed.
27/81 files processed.
28/81 files processed.
29/81 files processed.
30/81 files processed.
31/81 files processed.
```

32/81 files processed.  
33/81 files processed.  
34/81 files processed.  
35/81 files processed.  
36/81 files processed.  
37/81 files processed.  
38/81 files processed.  
39/81 files processed.  
40/81 files processed.  
41/81 files processed.  
42/81 files processed.  
43/81 files processed.  
44/81 files processed.  
45/81 files processed.  
46/81 files processed.  
47/81 files processed.  
48/81 files processed.  
49/81 files processed.  
50/81 files processed.  
51/81 files processed.  
52/81 files processed.  
53/81 files processed.  
54/81 files processed.  
55/81 files processed.  
56/81 files processed.  
57/81 files processed.  
58/81 files processed.  
59/81 files processed.  
60/81 files processed.  
61/81 files processed.  
62/81 files processed.  
63/81 files processed.  
64/81 files processed.  
65/81 files processed.  
66/81 files processed.  
67/81 files processed.  
68/81 files processed.  
69/81 files processed.  
70/81 files processed.  
71/81 files processed.  
72/81 files processed.  
73/81 files processed.  
74/81 files processed.  
75/81 files processed.  
76/81 files processed.  
77/81 files processed.  
78/81 files processed.  
79/81 files processed.

80/81 files processed.  
81/81 files processed.  
30 files found in data/log\_data  
1/30 files processed.  
2/30 files processed.  
3/30 files processed.  
4/30 files processed.  
5/30 files processed.  
6/30 files processed.  
7/30 files processed.  
8/30 files processed.  
9/30 files processed.  
10/30 files processed.  
11/30 files processed.  
12/30 files processed.  
13/30 files processed.  
14/30 files processed.  
15/30 files processed.  
16/30 files processed.  
17/30 files processed.  
18/30 files processed.  
19/30 files processed.  
20/30 files processed.  
21/30 files processed.  
22/30 files processed.  
23/30 files processed.  
24/30 files processed.  
25/30 files processed.  
26/30 files processed.  
27/30 files processed.  
28/30 files processed.  
29/30 files processed.  
30/30 files processed.

In [ ]:

In [ ]:

In [ ]: